

SEMINAR 5

SYSCALLS

AND A BIT OF X86 ASSEMBLY


PREVIOUSLY
WE USED A LOW-LEVEL
LANGUAGE BUT RELIED
ON GLIBC TO MANAGE EFFECTS

DEMO

BARE METAL

WONDERFUL
(ALBEIT UNSAFE)
DOS

PRIVILEGE RINGS



HOW SYSTEM CALL IS HANDLED

- > ENABLE PROTECTION OF USERSPACE CODE
- > REPLACE USER STACK WITH KERNEL STACK
 - > SAVE REGISTERS ON STACK
 - > PROCESS CALL
 - > RESTORE REGISTERS AND STACK
 - > DISABLE PROTECTION
 - > EXIT SYSTEM CALL

INT 80H

**OLDER MECHANISM
THAT WORKS ONLY WITH
32-BIT ADDRESSES**

SYSTEMENTER / SYSEXIT
OMITS CERTAIN CHECKS
TO WORK FASTER

HOWEVER

IT IS NOT

AMD-COMPATIBLE

SYSCALL/SYSRET
INTRODUCED BY AMD
SIMILAR TO SYSENTER
BUT USABLE

VSYSCALL

ALLOWS TO EXECUTE CERTAIN
SYSCALLS IN USERSPACE

VDSO

(VIRTUAL DYNAMIC SHARED OBJECT)

BETTER IMPLEMENTATION

OF VSYSCALL

PERFORMANCE COMPARISON

IMPLEMENTATION	TIME W/O KPTI (NS)	TIME W/ KPTI (NS)
INT 80H	317	498
SYSENTER	150	338
SYSCALL	103	278
VSYSCALL (EMUL)	496	692
VSYSCALL (NATIVE)	103	278
VDSO	37	37

A FEW WORDS ABOUT X86 ARCHITECTURE

> CISC

> 16 GENERAL-PURPOSE 64-BIT REGISTERS: **RAX, RBX, RCX, RDX, RBP, RSI, RDI, RSP, R8-R15**

> **RSP IS STACK POINTER**

> LINUX USES CALLING CONVENTION DEFINED IN **SYSTEM V AMD64 ABI**

SYSTEM V AMD64 ABI CALLING CONVENTION

- > ARGUMENTS PASSED IN **RDI, RSI, RDX, RCX, R8, R9**, REST ON STACK
 - > RETURN VALUE IN **RAX**
- > **RBX, RSP, RBP, R12-R15** ARE CALLEE-SAVED
 - > REST ARE CALLER-SAVED

LINUX SYSCALL CALLING CONVENTION

- > RAX FOR SYSCALL NUMBER
- > RDI, RSI, RDX, R10, R8, R9 FOR ARGUMENTS
 - > RAX, RDX FOR RETURN VALUE
 - > RCX AND R11 ARE DESTROYED

USEFUL LINKS

- > SYSCALL TABLE
- > X86-64 INSTRUCTION SET